

Deep Research Report: PINNsFormer in Scientific Machine Learning

Executive Summary

Physics-Informed Neural Networks (PINNs) have emerged as a powerful paradigm for solving partial differential equations (PDEs) by embedding physical laws into the loss function of neural networks. However, classical PINNs, which predominantly rely on Multilayer Perceptrons (MLPs), struggle with several failure modes, including spectral bias, gradient pathologies, and difficulties in capturing long-range temporal dependencies in stiff or chaotic systems. To address these limitations, the scientific machine learning community has increasingly turned to Transformer-based architectures. This report provides a comprehensive literature review and technical analysis of PINNsFormer and its major follow-up works, such as PhysicsFormer and Spectral PINNsFormer. We reverse-engineer the architectural patterns, mathematical formulations, and training methodologies that enable these models to outperform classical PINNs across various PDE and ODE systems.

Part 1: Historical Context

Classical PINNs and Their Limitations

The original PINN formulation, introduced by Raissi et al. [1](#), utilizes an MLP architecture to approximate the solution $u(x, t)$ of a PDE. The network takes spatial and temporal coordinates (x, t) as inputs and outputs the predicted solution. The physics is enforced through a composite loss function:

$$\mathcal{L} = \mathcal{L}_{res} + \lambda_{ic}\mathcal{L}_{ic} + \lambda_{bc}\mathcal{L}_{bc}$$

where $\mathcal{L}_{res}$ is the PDE residual loss computed via automatic differentiation, and $\mathcal{L}_{ic}$ and $\mathcal{L}_{bc}$ enforce initial and boundary conditions, respectively.

Despite their elegance, classical PINNs suffer from several well-documented failure modes:

- **Spectral Bias:** MLPs inherently prioritize learning low-frequency components of a function before high-frequency ones. Mathematically, this is explained by the Neural Tangent Kernel (NTK) theory, which shows that the eigenvalues of the NTK corresponding to high frequencies decay rapidly. This makes PINNs struggle with highly oscillatory or multi-scale PDEs [2](#).

- **Gradient Pathologies:** The composite loss function often leads to stiff optimization landscapes. The gradients of the residual loss can overwhelm the gradients of the boundary/initial condition losses, preventing the network from satisfying the constraints. This is a form of numerical stiffness in the gradient flow ³.
- **Multi-scale Dynamics and Stiffness:** In systems with widely varying time scales (e.g., chemical kinetics), the MLP struggles to capture both the fast transients and the slow evolution simultaneously.
- **Long-time Integration Errors:** Because classical PINNs treat time t merely as another continuous input coordinate, they fail to explicitly model the causal, sequential nature of time. This leads to error accumulation when integrating over long time horizons.

Part 2: PINNsFormer Core Architecture

To overcome the limitations of MLPs in capturing temporal dependencies, Zhao et al. proposed PINNsFormer ⁴. The core innovation is the adaptation of the Transformer architecture to continuous physical systems.

Architecture Breakdown

1. **Input Representation & Pseudo-Sequence Construction:** Unlike standard Transformers that process discrete tokens, PINNsFormer takes continuous coordinates (x, t) . To leverage attention, it constructs a "pseudo-sequence" by expanding the point-wise input into a sequence of historical time steps.
2. **Positional Encoding:** A linear embedding layer projects the (x, t) coordinates into a higher-dimensional latent space d_{model} .
3. **Transformer Encoder-Decoder:** The model utilizes a standard multi-head self-attention mechanism. The attention mechanism allows the model to weigh the importance of different historical time steps when predicting the current state, effectively capturing long-range temporal dependencies.
4. **Wavelet Activation Function:** Instead of standard activations like ReLU or Tanh, PINNsFormer introduces a Wavelet activation function: $\phi(x) = w_1 \sin(x) + w_2 \cos(x)$, where w_1, w_2 are learnable parameters. This anticipates Fourier decomposition, helping the network capture high-frequency components and mitigating spectral bias.

Mathematical Justification for Attention

In an MLP, the temporal evolution is implicitly learned through the global weights. In PINNsFormer, the self-attention mechanism explicitly computes a correlation matrix

between different time steps:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

This allows the network to dynamically route information from the initial conditions and past states to the current prediction, which is mathematically better suited for time-dependent PDEs than the static routing of an MLP.

Part 3: Physics Loss Construction

PINNsFormer modifies the standard point-wise PINN loss into a **sequential loss**.

Sequential Loss Formulation

Instead of evaluating the loss at isolated collocation points, PINNsFormer evaluates the loss over the constructed pseudo-sequences.

- **PDE Residual Loss:** $\mathcal{L}_{res} = \frac{1}{N} \sum_{i=1}^N \|\mathcal{F}(u_{\theta}(x_i, t_i))\|^2$
- **Initial Condition Loss:** $\mathcal{L}_{ic} = \frac{1}{N_{ic}} \sum_{i=1}^{N_{ic}} \|u_{\theta}(x_i, 0) - u_0(x_i)\|^2$
- **Boundary Condition Loss:** $\mathcal{L}_{bc} = \frac{1}{N_{bc}} \sum_{i=1}^{N_{bc}} \|\mathcal{B}(u_{\theta}(x_b, t))\|^2$

Advantages: The sequential loss forces the network to satisfy the PDE not just at a point, but along a temporal trajectory, improving the propagation of initial conditions.

Computational Cost: The self-attention mechanism scales as $O(L^2)$ where L is the sequence length, making it more computationally expensive than MLPs for very long sequences.

Part 4: PDE Families Studied

PDE	Type	Linear/Nonlinear	Dynamics	PINNsFormer Performance
Burgers	Parabolic	Nonlinear	Shock-dominated	Excellent; captures sharp gradients better than MLP.
Wave	Hyperbolic	Linear	Oscillatory	High accuracy; Wavelet

				activation helps with frequencies.
Allen-Cahn	Parabolic	Nonlinear	Multi-scale	Strong; handles phase transitions well.
Navier-Stokes	Fluid Dynamics	Nonlinear	Chaotic/Turbulent	Good for laminar/transitional; struggles with high Re turbulence.
Kuramoto-Sivashinsky	4th-order	Nonlinear	Chaotic	Improved over MLP, but long-time forecasting remains hard.

Why Transformers Help: For hyperbolic equations (Wave) and advection-dominated flows (Burgers), information propagates along characteristics. The attention mechanism can learn to align its attention weights with these characteristic curves, effectively tracking the flow of information.

Part 5: ODE Systems and Stiffness

Stiff vs. Non-Stiff ODEs

- **Non-stiff (e.g., Harmonic Oscillator, Lotka-Volterra):** PINNsFormer easily captures the oscillatory behavior, aided significantly by the Wavelet activation.
- **Stiff (e.g., Robertson Problem, Chemical Kinetics):** Stiff systems exhibit extreme time-scale separation (e.g., stiffness ratios $> 10^6$). The Robertson problem involves three chemical species reacting at vastly different rates.

Does PINNsFormer Solve Stiffness?

While PINNsFormer improves upon classical PINNs, **attention alone does not directly solve numerical stiffness**. The sequence modeling helps propagate information, but the gradient pathologies caused by the stiff PDE residuals still exist. State-of-the-art methods for stiff ODEs often require specialized techniques like Stiff-PINN (which uses quasi-steady-state approximations) or Reduced-PINN. PINNsFormer must be combined with adaptive loss weighting (like NTK) to handle severe stiffness.

Part 6: Transformer Design Variants

Model	Architecture	Key Innovation	Advantages
PINNsFormer 4	Encoder-Decoder	Pseudo-sequences, Wavelet activation	Strong temporal modeling, mitigates spectral bias.
PhysicsFormer 5	Cross-Attention	Efficient spatio-temporal embeddings	Faster training, highly efficient for Navier-Stokes.
Spectral PINNsFormer 6	Decoder-only	Fourier features, removes encoder	Reduces parameter count, explicitly solves spectral bias.

Spectral PINNsFormer (S-Pformer): Recent work 6 identified that the encoder in PINNsFormer is redundant for continuous coordinate inputs. By removing the encoder and injecting Fourier features directly into a decoder-only architecture, S-Pformer achieves higher accuracy with significantly fewer parameters.

Part 7: Mathematical Taxonomy

Where PINNsFormer Succeeds

- **1D and 2D Time-dependent PDEs:** Excellent performance on Burgers, Allen-Cahn, and low-Re Navier-Stokes.
- **Oscillatory Dynamics:** The Wavelet activation makes it highly effective for wave equations.

Where PINNsFormer Struggles

- **Extremely Stiff Systems:** Requires heavy hyperparameter tuning and NTK weighting.
- **High-Dimensional PDEs (>10D):** While better than MLPs, the sequence construction and attention mechanism suffer from the curse of dimensionality in terms of memory and compute.
- **Fully Developed 3D Turbulence:** Long-time forecasting of chaotic Navier-Stokes remains an open challenge for all PINN variants.

Part 8: Training Methodologies

Training Transformer-based PINNs requires careful optimization:

1. **Optimizers:** A two-stage approach is standard. **Adam** is used for initial exploration (e.g., 10,000 epochs) to navigate the complex loss landscape, followed by **L-BFGS** to exploit local convexity and achieve high precision.
2. **NTK Weighting:** To combat gradient pathologies, the loss weights λ_i are dynamically updated based on the trace of the Neural Tangent Kernel (NTK) for each loss component.
3. **Curriculum Learning / Adaptive Sampling:** For shock-dominated PDEs, collocation points are adaptively sampled in regions of high residual error.

Part 9: Code-Level Reverse Engineering

Based on the official PINNsFormer repository [4](#) :

- **Network Depth:** Typically $N = 1$ to $N = 3$ Transformer layers.
- **Hidden Dimensions:** $d_{model} = 32$, $d_{hidden} = 256$ or 512 .
- **Heads:** 2 to 4 attention heads.
- **Sequence Construction:** The `make_time_sequence` function duplicates the spatial coordinate and appends historical time steps (e.g., $t, t - \Delta t, t - 2\Delta t$) to create a sequence of length 5.

Pseudocode for Forward Pass:

Python

```
def forward(x, t):
    # x, t are pseudo-sequences of shape (Batch, Seq_Len, 1)
    src = concat(x, t)
    src_emb = Linear(src) # Project to d_model

    # Encoder
    e_out = src_emb
    for layer in encoder_layers:
        e_out = layer(e_out) # Self-attention + Wavelet FFN

    # Decoder
    d_out = src_emb
    for layer in decoder_layers:
        d_out = layer(d_out, e_out) # Cross-attention
```

```
# Output Head
pred = MLP_with_Wavelet(d_out)
return pred
```

Part 10: Research Gaps and Roadmap

Top Open Research Questions

1. How can we scale attention mechanisms to handle 3D fully developed turbulence without $O(L^2)$ memory explosion?
2. Can linear attention or state-space models (e.g., Mamba) replace standard attention for infinite-horizon time integration?
3. How do we theoretically guarantee the stability of Transformer-PINNs for stiff chemical kinetics?
4. Can we design a unified architecture that seamlessly transitions between operator learning (like FNO) and PINNsFormer?

Roadmap for Graduate Students

1. **Foundations:** Master classical PINNs, NTK theory, and the standard Transformer architecture.
2. **Implementation:** Replicate the S-Pformer [6](#) (decoder-only with Fourier features) as it is the most efficient current baseline.
3. **Contribution:** Focus on hybridizing Transformers with Neural Operators (e.g., Attention-based Neural Operators) or applying them to high-dimensional ($> 10D$) Hamilton-Jacobi-Bellman equations.

References

- [1] Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*.
- [2] Wang, S., Teng, Y., & Perdikaris, P. (2021). Understanding and mitigating gradient flow pathologies in physics-informed neural networks. *SIAM Journal on Scientific Computing*.
- [3] Ji, W., et al. (2021). Stiff-PINN: Physics-Informed Neural Network for Stiff Chemical Kinetics. *The Journal of Physical Chemistry A*.
- [4] Zhao, Z., Ding, X., & Prakash, B. A. (2023). PINNsFormer: A Transformer-Based Framework For Physics-Informed Neural Networks. *ICLR 2024*.

[5] Barman, B., Chatterjee, D., & Ray, R. K. (2026). PhysicsFormer: An Efficient and Fast Attention-Based Physics-Informed Neural Network for Solving Incompressible Navier-Stokes Equations.

[6] Arni, R. (2025). Physics-Informed Neural Networks with Fourier Features and Attention-Driven Decoding.